

Assignment 4: Kernels

April 16, 2021

0.0.1 Question 1

Proof To find vector w that minimizes $\sum_{n=1}^N \left(w \cdot x^{(n)} - y(n) \right)^2 + \lambda \|w\|^2$

We define X with $x^{(n)}$ on the n^{th} row and Y with $y^{(n)}$ on n^{th} row.

Therefore, learning objective becomes: $\|Xw - Y\|^2 + \lambda \|w\|^2$ We calculate gradient wrt w , $\nabla R(w)$

$$\nabla R(w) = 2X^T(Xw - Y) + 2\lambda w. \rightarrow (1)$$

For optimal $w \Rightarrow w^*$

$$\nabla R(w) = 0$$

by solving the equation we get

$$w^* = \left(X^T X + \lambda I \right)^{-1} X^T Y$$

We know that,

$$X^T X = \sum_{n=1}^N x^{(n)} x^{(n)T}$$

$$X^T Y = \sum_{n=1}^N x^{(n)} y^{(n)}$$

Resubstituting in equation (1), we get - $w^* = \left(\sum_{n=1}^N x^{(n)} x^{(n)T} + \lambda I \right)^{-1} \sum_{n=1}^N x^{(n)} y^{(n)}$

0.0.2 Question 2

Proof To find the vector w that minimizes

$$\sum_{n=1}^N \left(\omega \cdot \left(h(x^{(n)}) - y^{(n)} \right)^2 + \lambda \|w\|^2 \right).$$

where a basis expansion $h(x)$ is applied x

We define X with $x^{(n)}$ on the n^{th} row and H with $h(x)^{(n)}$ on the n^{th} row.

To find optimal w^*

$$w^* = \left(X^T X + \lambda I \right)^{-1} X^T y \rightarrow (1)$$

Substituting for basis expanded x we get

$$w^* = \left(H^T H + \lambda I \right)^{-1} H^T y \rightarrow (2)$$

$$H^T H = \sum_{n=1}^N h \left(x^{(n)} \right) \cdot h \left(x^{(n)} \right)^T \rightarrow (3)$$

$$H^T y = \sum_{n=1}^N h \left(x^{(n)} \right) \cdot y^{(n)} \rightarrow (4)$$

Substituting for $H^T H$ and $H^T y$ in (2)

$$\omega^* = \left(\sum_{n=1}^N h \left(x^{(n)} \right) \cdot h \left(x^{(n)} \right)^T + \lambda I \right)^{-1} \sum_{n=1}^N h \left(x^{(n)} \right) \cdot y^{(n)}$$

0.0.3 Question 3

Proof To derive an expression for $y^{\text{pred}} = w \cdot h(x^{\text{pred}})$ in terms of a kernel function k given $k(x, x') = h(x) \cdot h(x')$

From the derivation in Question 2 we get -

$$w^* = \left(H^T H + \lambda I \right)^{-1} H^T y \rightarrow (1)$$

Here, we define $h(x)^{(n)}$ on the n^{th} row.

We use a linear algebra trick for two matrices P & Q ,

$$(PQ + I)^{-1} P = P(QP + I)^{-1}$$

By applying this trick to equation (1), we will get -

$$w^* = H^T \left(HH^T + \lambda I \right)^{-1} y$$

Substituting HH^T with kernel matrix K , we take $\alpha = (K + \lambda I)^{-1} Y$.

$$w^* = H^T \alpha$$

The prediction is given by

$$y^{\text{pred}} = w^{*T} \cdot h(x^{\text{pred}})$$

Substituting value of w^* we get

$$\begin{aligned} y^{\text{pred}} &= (H^T \alpha)^T \cdot h(x^{\text{pred}}) \\ \Rightarrow y^{\text{pred}} &= \alpha^T H \cdot h(x^{\text{pred}}) \end{aligned}$$

Transforming it into summation form

$$y^{\text{pred}} = \sum_{n=1}^N \alpha_n k(x^{(n)}, x)$$

0.0.4 Question 4

Proof To find - A kernel function $k(x, x')$ where

$$k(x, x') = h(x) \cdot h(x')$$

$h(x)$ is the polynomial basis expansion

$$h(x) = [c_0, c_1x, c_2x^2, \dots, c_px^p]$$

$$cp = \sqrt{\binom{P}{p}}$$

We calculate $h(x')$ as -

$$h(x') = [c_0, c_1(x'), c_2(x'^2), \dots, cp(x'^p)]$$

Then, we take the dot product of h and $h(x')$ to evaluate $k(x, x')$ -

$$\begin{aligned} k(x, x') &= [c_0, c_1x, c_1x^2, \dots, c_px^p] \cdot [c_0, c_1x', c_2x^2, \dots, c_px'^p] \\ k(x, x') &= c_0^2 + c_1^2(x \cdot x') + c_2^2(x \cdot x')^2 + c_3^2(x \cdot x')^3 + \dots + c_p^2(x \cdot x')^p \end{aligned}$$

Using binomial expansion

$$c_0^2 + c_1^2(x \cdot x') + c_2^2(x \cdot x')^2 + c_3^2(x \cdot x')^3 + \dots + c_p^2(x \cdot x')^p = (1 + xx')^P$$

Hence our kernel matrix expression reduces to

$$k(x, x') = (1 + xx')^P$$

0.0.5 Question 5

```
[5]: def eval_basis_expanded_ridge(x,w,h):  
      expansion = h(x)  
      y = np.dot(w,expansion.T)  
      return y
```

0.0.6 Question 6


```
[6]: def train_basis_expanded_ridge(X,Y,,h = get_poly_expansion(3)):  
      H = h(X)  
      C = np.matmul(H.T,H) + *np.identity(H.shape[1])  
      w = np.linalg.solve(C,np.matmul(H.T,Y))  
      return w
```



0.0.7 Question 7


```
[ ]: P = [1,2,3,5,10]
      = 0.1
      for p in P:
          h = get_poly_expansion(p)
          w = train_basis_expanded_ridge(X_trn,Y_trn,,h)
          print("Weight vector = " + str(w) + " for P = " + str(p) + "\n")
```

P	Weight Vector
1	[1.00565302 0.12351259]
2	[1.55636445 -0.09905134 0.02105954]
3	[2.2585207 -0.4731082 0.09166343 -0.0074039]
5	[2.32031734e-01 1.70733532e+00 -7.50673263e-01 1.66018570e-01 -2.11820885e-02 1.49990203e-03]
10	[1.00207181e+00 3.45945567e-01 -6.77619200e-02 4.03404420e-02 -2.40833190e-02 7.29381906e-03 -1.30553266e-03 1.48669952e-04 -1.05198850e-05 3.68789523e-07 2.88162562e-09]

```
[8]: import matplotlib.pyplot as plt
      P = [1,2,3,5,10]
      = 0.1
      for p in P:
          h = get_poly_expansion(p)
          w = train_basis_expanded_ridge(X_trn,Y_trn,,h)
          x = np.linspace(0,15,1000,endpoint = True)
          fw = eval_basis_expanded_ridge(x,w,h)
          plt.scatter(x,fw,s=15,label = "Learned Function Output for P =" + str(p))
          plt.scatter(X_trn,Y_trn,s =15,label = "Training Data")
          plt.legend()
          plt.show()
```







0.0.8 Question 8

```
[9]: def get_poly_kernel(P):  
      def k(x, xp):  
          kernel_value = (1 + np.inner(x, xp)) ** P  
          return kernel_value  
      return k
```

0.0.9 Question 9

```
[ ]: x = 0.5
xp = 0.7
k = get_poly_kernel(5)
h = get_poly_expansion(5)
out1 = k(x,xp)
out2 = np.inner(h(x),h(xp))
print("output 1", out1)
print("output 2", out2)
```

Output 1(Kernel)	4.484033437500002
Output 2(Basis Expansion)	4.48403344

0.0.10 Question 10


```
[12]: def train_kernel_ridge(X,Y,,k):  
    # Initiate a Kernel matrix  
    K = np.zeros((X.shape[0],X.shape[0]))  
    # Fill values for kernel matrix using kernel function  
    for i in range(X.shape[0]):  
        for j in range(X.shape[0]):  
            K[i][j] = k(X[i],X[j])  
    # Get first part for np.linalg.solve  
    kli = K+np.identity(K.shape[0])  
    # Compute alpha  
    = np.linalg.solve(kli,Y)  
    return
```

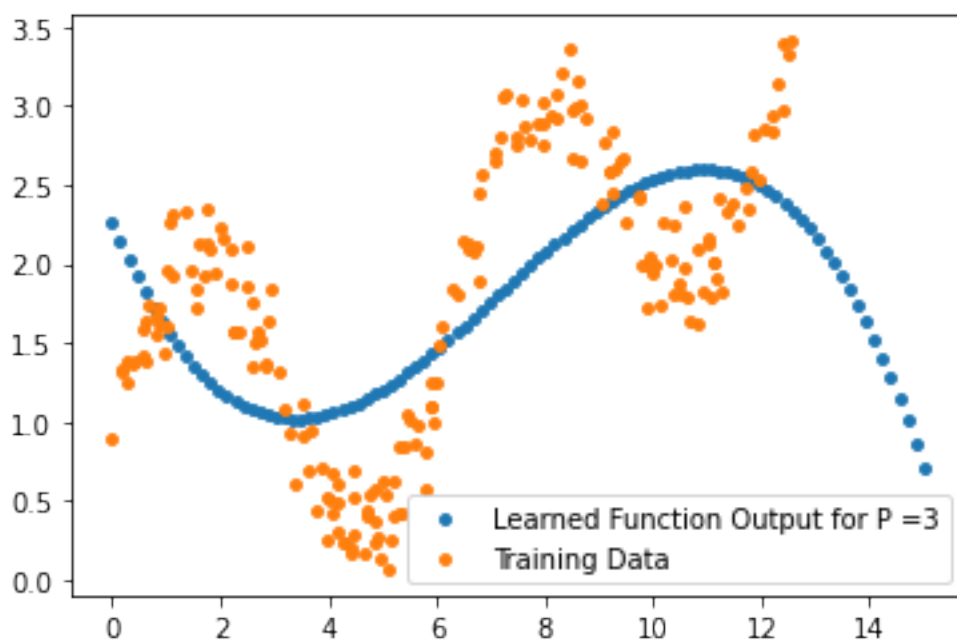

0.0.11 Question 11

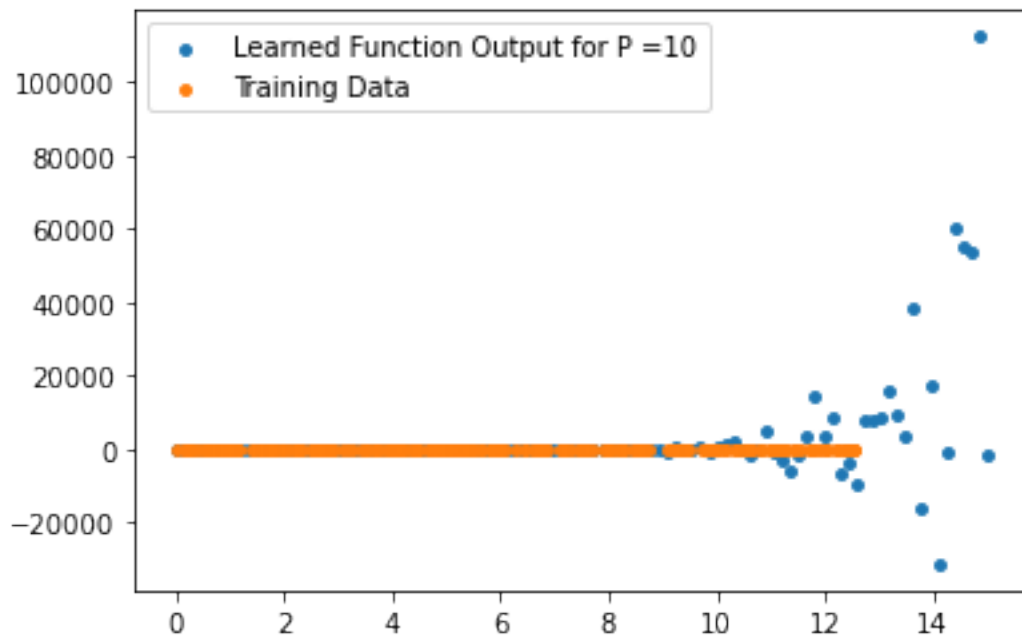

```
[13]: def eval_kernel_ridge(X_trn, x, , k):  
      # Evaluation of kernel ridge regression  
      sum_all = 0  
      for i in range(len(X_trn)):  
          sum_all += [i]*k(X_trn[i],x)  
      y = sum_all  
      return y
```

0.0.12 Question 12

```
[14]: import matplotlib.pyplot as plt
P = [1,2,3,5,10]
alpha = 0.1
for p in P:
    k = get_poly_kernel(p)
    model = train_kernel_ridge(X_trn,Y_trn,,k = get_poly_kernel(p))
    x = np.linspace(0,15,100)
    Y = []
    for i in range(len(x)):
        Y.append(eval_kernel_ridge(X_trn, x[i], , k))
    Y = np.array(Y)
    plt.scatter(x,Y,s=15,label = "Learned Function Output for P =" + str(p) )
    plt.scatter(X_trn,Y_trn,s =15,label = "Training Data")
    plt.legend()
    plt.show()
```







0.0.13 Question 13

The results of kernel ridge regression match that for basis expanded ridge regression for all values of P given in the question except for $P=10$ where the behaviour is a little erratic. The reason for the same can be found below:

According to Mercer's theorem, a Kernel K is only valid if it forms a positive definite matrix for all the datasets. When we check for all values of P , we observe that the kernel is not positive definite for P greater than 6 and hence it would be unwise to use kernel functions corresponding to those values of P (In our 5 cases, $P=10$ would not be valid as kernel matrix corresponding to that is not positive definite)

0.0.14 Question 14

```
[24]: stuff=np.load("data_real.npz")
      x_trn = stuff["x_trn"]
      y_trn = stuff["y_trn"]
      x_tst = stuff["x_tst"]

[ ]: from sklearn import svm
      from sklearn.model_selection import KFold
      from sklearn.metrics import hinge_loss

      c = [2,20,200]

      for penalty in c:
          hinge_losses = []

          kf = KFold(n_splits=5, shuffle=True, random_state=3815)
          for train_index, test_index in kf.split(x_trn):

              x_trn_5, x_tst_5 = x_trn[train_index], x_trn[test_index]
              y_trn_5, y_tst_5 = y_trn[train_index], y_trn[test_index]

              clf = svm.SVC(kernel = 'linear', C = 1/penalty)
              clf.fit(x_trn_5,y_trn_5)

              y_pred = clf.decision_function(x_tst_5)
              hinge_losses.append(hinge_loss(y_tst_5,y_pred))

      mean_hinge_loss = (sum(hinge_losses)/5)

      print("The mean hinge loss with 5-Fold cross validation using hinge loss_
→with lambda = " + str(penalty) + " is: ", mean_hinge_loss)
      print("*****")
```

HINGE LOSS FOR LINEAR KERNEL	
Lambda	Mean hinge loss
2	0.0392
20	0.0473
200	0.0827

0.0.15 Question 15

```
[ ]: from sklearn import svm
from sklearn.model_selection import KFold
from sklearn.metrics import hinge_loss

c = [2,20,200]
gamma = [1,0.01,0.001]
for penalty in c:
    for g in gamma:

        hinge_losses = []

        kf = KFold(n_splits=5, shuffle=True, random_state=3815)
        for train_index, test_index in kf.split(x_trn):

            x_trn_5, x_tst_5 = x_trn[train_index], x_trn[test_index]
            y_trn_5, y_tst_5 = y_trn[train_index], y_trn[test_index]

            clf = svm.SVC(kernel = 'poly', degree = 3, C = 1/penalty, gamma = 1,
→coef0 = g)
            clf.fit(x_trn_5,y_trn_5)
            y_pred = clf.decision_function(x_tst_5)
            hinge_losses.append(hinge_loss(y_tst_5,y_pred))

        mean_hinge_loss = (sum(hinge_losses)/5)
        print("The mean hinge loss with 5-Fold cross validation using hinge loss,
→with lambda = " + str(penalty) + " and gamma = " + str(g) + " is: " ,
→mean_hinge_loss)

        print("*****")
```

Mean Hinge Loss

Lambda, Gamma	1	0.01	0.001
2	0.0251	0.0684	0.0629
20	0.0509	0.0453	0.0365
200	0.0462	0.0545	0.0582

0.0.16 Question 16

```
[ ]: from sklearn import svm
from sklearn.model_selection import KFold
from sklearn.metrics import hinge_loss

c = [2,20,200]
gamma = [1,0.01,0.001]
for penalty in c:
    for g in gamma:

        hinge_losses = []

        kf = KFold(n_splits=5, shuffle=True, random_state=3815)
        for train_index, test_index in kf.split(x_trn):

            x_trn_5, x_tst_5 = x_trn[train_index], x_trn[test_index]
            y_trn_5, y_tst_5 = y_trn[train_index], y_trn[test_index]

            clf = svm.SVC(kernel = 'poly', degree = 5, C = 1/penalty, gamma = 1,
→coef0 = g)
            clf.fit(x_trn_5,y_trn_5)
            y_pred = clf.decision_function(x_tst_5)
            hinge_losses.append(hinge_loss(y_tst_5,y_pred))

        mean_hinge_loss = (sum(hinge_losses)/5)

        print("The mean hinge loss with 5-Fold cross validation using hinge loss,
→with lambda = " + str(penalty) + " and gamma = " + str(g) + " is: " ,
→mean_hinge_loss)
```

Mean Hinge Loss

Lambda, Gamma	1	0.01	0.001
2	0.0490	0.3603	0.3183
20	0.4674	1.0210	0.5283
200	0.0705	0.4236	0.4310

0.0.17 Question 17

```
[ ]: from sklearn import svm
from sklearn.model_selection import KFold
from sklearn.metrics import hinge_loss

c = [2,20,200]
gamma = [1,0.01,0.001]
for penalty in c:
    for g in gamma:
        hinge_losses = []
        errors = []

        kf = KFold(n_splits=5, shuffle=True, random_state=3815)
        for train_index, test_index in kf.split(x_trn):

            x_trn_5, x_tst_5 = x_trn[train_index], x_trn[test_index]
            y_trn_5, y_tst_5 = y_trn[train_index], y_trn[test_index]

            clf = svm.SVC(kernel = 'rbf', C = 1/penalty, gamma = g)
            clf.fit(x_trn_5,y_trn_5)

            y_pred = clf.decision_function(x_tst_5)
            hinge_losses.append(hinge_loss(y_tst_5,y_pred))

            y_pred2 = clf.predict(x_tst_5)
            accuracy = (sum(y_pred2 == y_tst_5))/len(y_tst_5)
            errors.append(1 - accuracy)

        mean_hinge_loss = (sum(hinge_losses)/5)
        mean_error = (sum(errors))/5

#         print("The mean hinge loss with 5-Fold cross validation using hinge_
→loss with lambda = " + str(penalty) + " and gamma = " + str(g) + " is: " ,
→mean_hinge_loss)
#         print("*****")
```

Mean Hinge Loss for RBF kernel

Lambda, Gamma	1	0.01	0.001
2	0.2476	0.0568	0.1605
20	0.7811	0.2029	0.6531
200	0.8783	0.7826	0.8655

Generalization Error for RBF Kernel			
Lambda, Gamma	1	0.01	0.001
2	0.0130	0.0102	0.0277
20	0.4445	0.0247	0.3148
200	0.4445	0.445	0.445

0.0.18 Question 18

1. We chose Polynomial kernel with following hyperparameters:

$$\gamma = 1, \lambda = 2$$

For this polynomial kernel, the hinge loss that we observe in the K-Fold cross validation data is minimum amongst all the other methods that we tried

2. Estimated Generalization Error is: 0.006%
3. Observed Generalization Error is: 0%